

A proof-governed autonomy system for career conversion in embodied creative research and development

Alan N. Pham^{1,*}

¹AO Labs; institutional affiliation and coauthor list to be confirmed before submission

*Correspondence: Alan Pham

Autonomous agents are increasingly able to plan, monitor, and execute tool-mediated workflows, but a career objective is normally treated as motivation rather than as a measurable control problem. We report a proof-governed system that converts a target role, Walt Disney Imagineering Research and Development mechanical Imagineering in Glendale, into a logged state machine with role-fit dimensions, guardrails, daily cycles, evidence counters, weekly manuscript updates, and an autonomous reviewer. The reviewer uses GPT-5.5 and a whole-public-AO-Labs evidence graph, including the CV, Sarrus, FluxCell, Relay, Relay Live, Ocean, other public project surfaces, the WDI proof packet, the active WDI listing, and Disney Research context. It can critique evidence autonomously but cannot send outreach, claim relationships, submit applications, or use private email/contact data without approval. At the current evidence lock on 6 May 2026, the system tracked six role-fit dimensions, four portfolio anchors, seven proof logs, one daily cycle, ten AI reviews, zero outreach events, five reviewer-ready artifacts, and twenty journal entries. The public dashboard is now report-first: the primary section is the reviewer-generated critique of Alan's current evidence, followed by reviewer-system status and the manuscript artifact. The live fit score was 90 of 100 with a strong-and-visible confidence label, while the autonomous reviewer scored the public packet 88 of 100 after the system published a Sarrus mechanism proof page, added an HTML-visible quantitative mechanism table, and inserted role-calibration boundaries for the active Glendale mechanical-design rung versus the Principal R&D Imagineer north-star profile. The current bottleneck is no longer absence of a mechanism artifact; it is review intelligence and level calibration: the strongest public evidence supports active WDI mechanical-design fit, while principal-level leadership evidence and an active principal posting remain unverified. The contribution is therefore not a claim that the system has secured a job. It is a falsifiable protocol for making an otherwise ambiguous career-conversion process auditable,

²⁹ **adaptive, and constrained by truthful evidence.**

1 Introduction

Tool-using language-model systems can now reason, call external services, write artifacts, and update deployed software (1–4). These capabilities create an opportunity to treat long-horizon personal objectives as monitored systems rather than as episodic intentions, in the older spirit of autonomic systems and closed-loop discovery workflows (5, 6). A career target is especially difficult because the desired outcome depends on evidence quality, technical fit, social trust, timing, applications, and human judgment. Without a state model, repeated effort can feel intense while remaining scientifically uninterpretable.

This paper describes an early deployment of an autonomous career-evidence system aimed at WDI Research and Development mechanical Imagineering roles. The system uses the active WDI Research and Development Imagineer - Mechanical Design Engineer role as the live rung and the Principal R&D Imagineer - Mechanical Engineer profile as the north-star target (7). The positioning line is deliberately stable: mechanical PhD, soft robotics, creative prototyping, and AI-assisted tools for human-facing physical experiences.

The design borrows from proof-governed revenue experimentation: define a measurable state, choose the current bottleneck, constrain allowed actions, record interventions, and update a paper from real evidence rather than from aspiration. The system is not allowed to invent credentials, inflate relationships, send spam, or apply on behalf of the operator. It can monitor state, update dashboards, maintain a paper, identify the next constraint, run an evidence-only AI review, and create or improve public artifacts that are grounded in existing work.

2 System architecture

The deployment consists of a public dashboard, a Railway-hosted backend, a JSON runtime state file mounted on a persistent volume, and scheduled or operator-triggered automation. The backend exposes health, operations-check, journal, weekly-paper, event, daily-cycle, and AI-review endpoints. The dashboard is intentionally minimal and report-first. Its three sections are: a reviewer-generated readout on Alan’s current public evidence, a compact status surface for the AI reviewer and role-fit dimensions, and the paper/PDF artifact. The dashboard does not treat the reviewer as the main subject; it treats the reviewer as an instrument that generates feedback for human judgment.

Table 1. State variables at the current evidence lock. Values are reported from the live system after deployment of the whole-public-AO-Labs GPT-5.5 reviewer on 6 May 2026.

Variable	Recorded value
Target company and location	Walt Disney Imagineering R&D, Glendale, California
Live rung	WDI Research and Development Imagineer - Mechanical Design Engineer
North star	Principal R&D Imagineer - Mechanical Engineer
Fit score	90 of 100 with strong-and-visible confidence
Dominant bottleneck	Review intelligence: repeatable critique, source coverage, and approved escalation
Evidence counters	Seven proof logs; one daily cycle; ten AI reviews; zero outreach events; four portfolio anchors; five reviewer-ready artifacts; twenty journal entries
Active experiment	Autonomous AI reviewer v0
Reviewer source graph	Twenty-five live sources in the latest production run, spanning public AO Labs projects and WDI/Disney Research context
Human approval boundary	Applications, sensitive outreach, referrals, claims involving real people, and any private contact/email data

The reviewer is a separate loop rather than a human-contact substitute. It builds a source packet from the live state, the target listing, the proof packet, the paper PDF, the CV, Sarrus, FluxCell, Relay, Relay Live, Ocean, Talk, Nerve, Duet, Violin, Yum, Lily, AO Labs home, Disney Research roster, and a Disney Research bipedal-character paper page. It also includes a compact identity profile summarizing the operator as a mechanical engineering PhD candidate building soft robotic materials, reconfigurable mechanisms, morphing interfaces, research papers, public project surfaces, and autonomous systems. The model returns a verdict, score, top issue, evidence gaps, packet edits, and next actions as JSON. If the model fails or times out, a deterministic fallback reviewer produces a conservative critique and records the fallback reason.

3 Decision policy

The policy scores six role-fit dimensions: mechanical depth, creative prototyping, human-facing physical experience, review intelligence, Glendale application packet, and the paper/system itself. Each dimension now reports what it measures, its score basis, and the next evidence that would move the signal. The score basis is intentionally mechanical: a base prior from the state file, plus logged event impact, portfolio evidence tags, and daily-cycle evidence, capped at 100. The next action is

chosen from the weakest current signal. In the current lock, the weakest signal is review intelligence. 73
 The selected intervention is to run the autonomous reviewer and route its highest-leverage critique into 74
 a concrete public artifact improvement. 75

This policy is intentionally conservative. It can make public artifacts, update the dashboard, refresh 76
 the manuscript, run autonomous critique, and record evidence. It cannot manufacture the social proof 77
 that a principal-level role requires. If an action requires a real person's attention or reputation, the 78
 system must stop at preparation and request approval before sending. The reviewer therefore becomes 79
 a repeatable calibration instrument, not a license to automate human outreach. 80

4 Results 81

The current deployment is functional and has moved beyond the initial dashboard-only state. The 82
 dashboard and backend are live, the weekly paper endpoint can publish a snapshot, the site exposes a 83
 PDF manuscript, and the AI reviewer has produced live production critiques from the public evidence 84
 graph. The latest reviewer run used GPT-5.5, consumed 25 sources, returned a reviewer score of 88 of 85
 100, and identified a narrower dominant issue: the strongest public evidence now supports the active 86
 Glendale WDI Research and Development Imagineer - Mechanical Design Engineer rung, while the 87
 Principal R&D Imagineer north-star remains unverified and under-evidenced. 88

The largest artifact change in this cycle was not a change in language; it was a change in inspectable 89
 proof. The system published a reviewer-facing Sarrus mechanism page with mechanism geometry, 90
 compliant-link and flexural-joint path, pressure-driven expansion, module assembly, prototype state, 91
 visual force/stiffness evidence, and show-value translation. It then added an HTML-visible quantitative 92
 table with pressure range, expansion ratio, estimated leg angle, output force, local stiffness, hysteresis 93
 status, planar dome height, double-dome resolution, and prototype limitation. The proof packet was 94
 updated to make this table the dominant evidence path and to preserve the current data boundary: raw 95
 digitized force, stiffness, and hysteresis curves with uncertainty remain a next proof item. 96

After the latest interface revision, the live site no longer leads with generic status tiles. It leads with 97
 the reviewer report: verdict, reviewer score, summary, top issue, best existing evidence, evidence 98
 gaps, packet edits, and the selected next action. The reviewer-system section remains visible, but its 99
 purpose is operational transparency: model, source count, review count, approval boundary, fit score, 100
 bottleneck, evidence counters, and role-fit score explanations. This distinction matters because the 101

Table 2. Latest autonomous reviewer result. The reviewer is a GPT-5.5 evidence critique over the public AO Labs graph, not a human recommendation.

Field	Recorded value
Model	GPT-5.5 via the Responses API
Scope	Whole public AO Labs graph plus WDI role and Disney Research context
Source count	25 sources in the latest production run
Reviewer score	88 of 100
Verdict	Strong active-rung fit, not yet principal-level proof
Top issue	Level calibration: the strongest evidence supports the active Glendale Mechanical Design Engineer rung, while the principal north-star remains unverified and under-evidenced
Next action	Keep the role-verification banner above the fold

useful artifact for the operator is the critique, not a status page about the critique generator.

The system-level fit score increased to 90 of 100, but this should be interpreted as improved instrumentation and artifact readiness, not as an external hiring probability. The reviewer result is deliberately stricter than the system score. It confirms that the overall direction is credible for WDI R&D mechanical design while rejecting the idea that the current public packet is already principal-level proof. The proof packet now includes both a principal-scope evidence boundary and an above-fold role-verification banner: the active target is the Glendale mechanical-design rung; the principal title remains a north-star profile until an active principal posting and principal-scope leadership evidence are independently verified.

5 Limitations and guardrails

The system cannot guarantee an offer. A job decision depends on timing, hiring needs, competition, interview performance, portfolio interpretation, and institutional context. The useful claim is narrower: the system can make progress legible, choose the next bottleneck, prevent fake progress, and turn repeated effort into evidence. It should be evaluated by proof velocity, reviewer-ready artifacts, quality of critique paths, application readiness, interviews, and eventual conversion.

The guardrails are part of the result. The system forbids fabricated credentials, fake outreach, invented Disney relationships, automated applications, and claims that are not supported by public or logged evidence. Private email/contact data is excluded from the reviewer unless explicitly approved. These

Table 3. Role-fit dimensions after the current autonomous reviewer cycle. Scores are internal operating metrics, not external hiring probabilities; the live dashboard exposes the measurement target, score basis, and next evidence for each row.

Dimension	Score	Current interpretation
Mechanical depth	100	Strong internal score from accumulated proof/system events; Sarrus now has a public mechanism page and quantitative evidence table, with raw digitized curves still pending.
Creative prototyping	82	Strongest signal; should remain attached to concrete physical demonstrations.
Human-facing physical experience	100	Strong internal score after the Sarrus object-manipulation and motion translation were made reviewer-visible.
Review intelligence	55	Current bottleneck; requires repeatable critique, source coverage, and optional approved human escalation.
Glendale packet	100	Strong internal readiness score after quantitative Sarrus evidence, role verification, and principal-scope boundaries were added.
Autonomous career system	100	Backend, dashboard, durable state, reviewer, journal, and paper loop are live.

constraints reduce short-term aggressiveness, but they preserve interpretability. If the system ever reports strong progress, that progress should be traceable to artifacts, logs, AI reviews, human approvals, applications, or outcomes.

6 Availability

The live dashboard is available at <https://imagineer.aolabs.io/>, with a fallback at <https://aolabs.io/imagineer/>. The live backend exposes <https://imagineer.aolabs.io/api/imagineer/ops-check>, <https://imagineer.aolabs.io/api/imagineer/weekly-paper>, and <https://imagineer.aolabs.io/api/imagineer/ai-review>. This manuscript is a public artifact and should be refreshed from logged evidence as the system accumulates real proof.

References

1. S. Yao *et al.*, ReAct: Synergizing reasoning and acting in language models. *International Conference on Learning Representations* (2023).

-
- 132 2. T. Schick *et al.*, Toolformer: Language models can teach themselves to use tools. *arXiv preprint*
133 *arXiv:2302.04761* (2023).
 - 134 3. Q. Wu *et al.*, AutoGen: Enabling next-gen LLM applications via multi-agent conversation
135 framework. *arXiv preprint arXiv:2308.08155* (2023).
 - 136 4. L. Wang *et al.*, A survey on large language model based autonomous agents. *arXiv preprint*
137 *arXiv:2308.11432* (2023).
 - 138 5. J. O. Kephart, D. M. Chess, The vision of autonomic computing. *Computer* **36**, 41–50 (2003).
 - 139 6. A. G. Kusne *et al.*, On-the-fly closed-loop materials discovery via Bayesian active learning. *Nature*
140 *Communications* **11**, 5966 (2020).
 - 141 7. The Walt Disney Company, *WDI Research and Development Imagineer - Mechanical Design*
142 *Engineer* (2026; [https://jobs.disneycareers.com/job/glendale/wdi-research-](https://jobs.disneycareers.com/job/glendale/wdi-research-and-development-imagineer-mechanical-design-engineer/391/93733641696)
143 [and-development-imagineer-mechanical-design-engineer/391/93733641696](https://jobs.disneycareers.com/job/glendale/wdi-research-and-development-imagineer-mechanical-design-engineer/391/93733641696)).